

Income_Qualification

May 4, 2021

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import collections
from collections import Counter
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score
from sklearn.model_selection import cross_val_score

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: # Loading the dataset
df_income_train = pd.read_csv("train.csv")
df_income_test = pd.read_csv("test.csv")
```

```
[3]: df_income_train.head()
```

```
[3]:
```

	Id	v2a1	hacdor	rooms	hacapo	v14a	refrig	v18q	v18q1	\
0	ID_279628684	190000.0	0	3	0	1	1	0	NaN	
1	ID_f29eb3ddd	135000.0	0	4	0	1	1	1	1.0	
2	ID_68de51c94	NaN	0	8	0	1	1	0	NaN	
3	ID_d671db89c	180000.0	0	5	0	1	1	1	1.0	
4	ID_d56d6f5f5	180000.0	0	5	0	1	1	1	1.0	

	r4h1	...	SQBescolari	SQBage	SQBhogar_total	SQBedjefe	SQBhogar_nin	\
0	0	...	100	1849	1	100	0	
1	0	...	144	4489	1	144	0	
2	0	...	121	8464	1	0	0	
3	0	...	81	289	16	121	4	
4	0	...	121	1369	16	121	4	

	SQBovercrowding	SQBdependency	SQBmeaned	agesq	Target
0	1.000000	0.0	100.0	1849	4
1	1.000000	64.0	144.0	4489	4
2	0.250000	64.0	121.0	8464	4
3	1.777778	1.0	121.0	289	4
4	1.777778	1.0	121.0	1369	4

[5 rows x 143 columns]

```
[4]: df_income_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9557 entries, 0 to 9556
Columns: 143 entries, Id to Target
dtypes: float64(8), int64(130), object(5)
memory usage: 10.4+ MB
```

```
[5]: df_income_test.head()
```

```
[5]:
```

	Id	v2a1	hacdor	rooms	hacapo	v14a	refrig	v18q	v18q1	\
0	ID_2f6873615	NaN	0	5	0	1	1	0	NaN	
1	ID_1c78846d2	NaN	0	5	0	1	1	0	NaN	
2	ID_e5442cf6a	NaN	0	5	0	1	1	0	NaN	
3	ID_a8db26a79	NaN	0	14	0	1	1	1	1.0	
4	ID_a62966799	175000.0	0	4	0	1	1	1	1.0	

	r4h1	...	age	SQBescolari	SQBage	SQBhogar_total	SQBedjefe	\
0	1	...	4	0	16	9	0	
1	1	...	41	256	1681	9	0	
2	1	...	41	289	1681	9	0	
3	0	...	59	256	3481	1	256	
4	0	...	18	121	324	1	0	

	SQBhogar_nin	SQBovercrowding	SQBdependency	SQBmeaned	agesq
0	1	2.25	0.25	272.25	16
1	1	2.25	0.25	272.25	1681
2	1	2.25	0.25	272.25	1681
3	0	1.00	0.00	256.00	3481
4	1	0.25	64.00	NaN	324

[5 rows x 142 columns]

```
[6]: df_income_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23856 entries, 0 to 23855
Columns: 142 entries, Id to agesq
dtypes: float64(8), int64(129), object(5)
```

memory usage: 25.8+ MB

```
[7]: #List the columns for different datatypes:
print('Integer Type: ')
print(df_income_train.select_dtypes(np.int64).columns)
print('\n')
print('Float Type: ')
print(df_income_train.select_dtypes(np.float64).columns)
print('\n')
print('Object Type: ')
print(df_income_train.select_dtypes(np.object).columns)
```

Integer Type:

```
Index(['hacdor', 'rooms', 'hacapo', 'v14a', 'refrig', 'v18q', 'r4h1', 'r4h2',
      'r4h3', 'r4m1',
      ...,
      'area1', 'area2', 'age', 'SQBescolari', 'SQBage', 'SQBhogar_total',
      'SQBedjefe', 'SQBhogar_nin', 'agesq', 'Target'],
      dtype='object', length=130)
```

Float Type:

```
Index(['v2a1', 'v18q1', 'rez_esc', 'meaneduc', 'overcrowding',
      'SQBovercrowding', 'SQBdependency', 'SQBmeaned'],
      dtype='object')
```

Object Type:

```
Index(['Id', 'idhogar', 'dependency', 'edjefe', 'edjefa'], dtype='object')
```

```
[8]: df_income_train.select_dtypes('int64').head()
```

```
[8]:   hacdor  rooms  hacapo  v14a  refrig  v18q  r4h1  r4h2  r4h3  r4m1  ...  \
0        0      3        0      1        1      0      0      1      1      0  ...
1        0      4        0      1        1      1      0      1      1      0  ...
2        0      8        0      1        1      0      0      0      0      0  ...
3        0      5        0      1        1      1      0      2      2      1  ...
4        0      5        0      1        1      1      0      2      2      1  ...

   area1  area2  age  SQBescolari  SQBage  SQBhogar_total  SQBedjefe  \
0        1      0  43          100    1849              1         100
1        1      0  67          144   4489              1          144
2        1      0  92          121   8464              1           0
3        1      0  17           81    289             16          121
4        1      0  37          121   1369             16          121

   SQBhogar_nin  agesq  Target
```

0	0	1849	4
1	0	4489	4
2	0	8464	4
3	4	289	4
4	4	1369	4

[5 rows x 130 columns]

```
[9]: for column in df_income_train:
      if column not in df_income_test:
          print('The output variable is', column)
```

The output variable is Target

```
[10]: print(df_income_train['Target'].value_counts())
```

```
4    5996
2    1597
3    1209
1     755
Name: Target, dtype: int64
```

```
[11]: # Finding columns with null values
null_counts=df_income_train.select_dtypes('int64').isnull().sum()
null_counts[null_counts > 0]
```

```
[11]: Series([], dtype: int64)
```

```
[12]: df_income_train.select_dtypes('float64').head()
```

```
[12]:
```

	v2a1	v18q1	rez_esc	meaneduc	overcrowding	SQBovercrowding	\
0	190000.0	NaN	NaN	10.0	1.000000	1.000000	
1	135000.0	1.0	NaN	12.0	1.000000	1.000000	
2	NaN	NaN	NaN	11.0	0.500000	0.250000	
3	180000.0	1.0	1.0	11.0	1.333333	1.777778	
4	180000.0	1.0	NaN	11.0	1.333333	1.777778	

	SQBdependency	SQBmeaned
0	0.0	100.0
1	64.0	144.0
2	64.0	121.0
3	1.0	121.0
4	1.0	121.0

```
[13]: # Finding columns with null values
null_counts=df_income_train.select_dtypes('float64').isnull().sum()
null_counts[null_counts > 0]
```

```
[13]: v2a1          6860
      v18q1        7342
      rez_esc      7928
      meaneduc      5
      SQBmeaned     5
      dtype: int64
```

```
[14]: df_income_train.select_dtypes('object').head()
```

```
[14]:
```

	Id	idhogar	dependency	edjefe	edjefa
0	ID_279628684	21eb7fcc1	no	10	no
1	ID_f29eb3ddd	0e5d7a658	8	12	no
2	ID_68de51c94	2c7317ea8	8	no	11
3	ID_d671db89c	2b58d945f	yes	11	no
4	ID_d56d6f5f5	2b58d945f	yes	11	no

```
[15]: #Find columns with null values
null_counts=df_income_train.select_dtypes('object').isnull().sum()
null_counts[null_counts > 0]
```

```
[15]: Series([], dtype: int64)
```

```
[16]: # With out Null values treatment we cannot get the correct answers
```

0.1 Define Variable Categories

There are several different categories of variables:

1. **Squared Variables:** derived from squaring variables in the data
2. **Id variables:** identifies the data and should not be used as features
3. **Household variables**
 - Boolean: Yes or No
 - Ordered Discrete: Integers with an ordering
 - Continuous numeric
4. **Individual Variables:** these are characteristics of each individual rather than the household
 - Boolean: Yes or No (0 or 1)
 - Ordered Discrete: Integers with an ordering

```
[19]: # dependency column
```

```
[17]: mapping={'yes':1,'no':0}

for df in [df_income_train, df_income_test]:
```

```
df['dependency'] =df['dependency'].replace(mapping).astype(np.float64)
df['edjefe'] =df['edjefe'].replace(mapping).astype(np.float64)
df['edjefa'] =df['edjefa'].replace(mapping).astype(np.float64)
```

```
[18]: df_income_train[['dependency','edjefe','edjefa']]
```

```
[18]:
```

	dependency	edjefe	edjefa
0	0.00	10.0	0.0
1	8.00	12.0	0.0
2	8.00	0.0	11.0
3	1.00	11.0	0.0
4	1.00	11.0	0.0
...
9552	0.25	9.0	0.0
9553	0.25	9.0	0.0
9554	0.25	9.0	0.0
9555	0.25	9.0	0.0
9556	0.25	9.0	0.0

[9557 rows x 3 columns]

```
[19]: df_income_train[['dependency','edjefe','edjefa']].describe()
```

```
[19]:
```

	dependency	edjefe	edjefa
count	9557.000000	9557.000000	9557.000000
mean	1.149550	5.096788	2.896830
std	1.605993	5.246513	4.612056
min	0.000000	0.000000	0.000000
25%	0.333333	0.000000	0.000000
50%	0.666667	6.000000	0.000000
75%	1.333333	9.000000	6.000000
max	8.000000	21.000000	21.000000

```
[23]: df_income_test[['dependency','edjefe','edjefa']].describe()
```

```
[23]:
```

	dependency	edjefe	edjefa
count	23856.000000	23856.000000	23856.000000
mean	1.181327	5.199824	2.800176
std	1.666209	5.200980	4.603592
min	0.000000	0.000000	0.000000
25%	0.333333	0.000000	0.000000
50%	0.666667	6.000000	0.000000
75%	1.333333	9.000000	6.000000
max	8.000000	21.000000	21.000000

```
[24]: # v2a1 column
```

```
[25]: data = df_income_train[df_income_train['v2a1'].isnull()].head()

columns=['tipovivi1','tipovivi2','tipovivi3','tipovivi4','tipovivi5']
data[columns]
```

```
[25]:
```

	tipovivi1	tipovivi2	tipovivi3	tipovivi4	tipovivi5
2	1	0	0	0	0
13	1	0	0	0	0
14	1	0	0	0	0
26	1	0	0	0	0
32	1	0	0	0	0


```
[26]: # Variables indicating home ownership
own_variables = [x for x in df_income_train if x.startswith('tipo')]

# Plot of the home ownership variables for home missing rent payments
df_income_train.loc[df_income_train['v2a1'].isnull(), own_variables].sum().plot.
    ↳ bar(figsize = (10, 8),

                                                    color = '
    ↳ 'red',

                                                    edgecolor = 'k',
    ↳ linewidth = 2);
plt.xticks([0, 1, 2, 3, 4],
            ['Owns and Paid Off', 'Owns and Paying', 'Rented', 'Precarious',
    ↳ 'Other'],
            rotation = 20)
plt.title('Home Ownership Status for Households Missing Rent Payments', size =
    ↳ 18);
```

Home Ownership Status for Households Missing Rent Payments



```
[27]: for df in [df_income_train, df_income_test]:
        df['v2a1'].fillna(value=0, inplace=True)

df_income_train[['v2a1']].isnull().sum()
```

```
[27]: v2a1    0
      dtype: int64
```

```
[28]: df_income_test[['v2a1']].isnull().sum()
```

```
[28]: v2a1    0
      dtype: int64
```


```
[29]: # v18q1 Column
```

```
[30]: heads = df_income_train.loc[df_income_train['parentesco1'] == 1].copy()
heads.groupby('v18q1')['v18q1'].apply(lambda x: x.isnull().sum())
```



```
[30]: v18q1
0    2318
1         0
Name: v18q1, dtype: int64
```

```
[31]: plt.figure(figsize = (8, 6))
col='v18q1'
df_income_train[col].value_counts().sort_index().plot.bar(color = 'blue',
                                                            edgecolor = 'k',
                                                            linewidth = 2)
plt.xlabel(f'{col}'); plt.title(f'{col} Value Counts'); plt.ylabel('Count')
plt.show();
```



```
[32]: for df in [df_income_train, df_income_test]:
        df['v18q1'].fillna(value=0, inplace=True)

df_income_train[['v18q1']].isnull().sum()
```

```
[32]: v18q1    0
dtype: int64
```

```
[33]: df_income_test[['v18q1']].isnull().sum()
```

```
[33]: v18q1    0  
      dtype: int64
```

```
[34]: # rez_esc column
```

```
[35]: # Checking for no null values  
df_income_train[df_income_train['rez_esc'].notnull()]['age'].describe()
```

```
[35]: count    1629.000000  
      mean      12.258441  
      std       3.218325  
      min       7.000000  
      25%       9.000000  
      50%      12.000000  
      75%      15.000000  
      max      17.000000  
      Name: age, dtype: float64
```

```
[36]: df_income_train.loc[df_income_train['rez_esc'].isnull()]['age'].describe()
```

```
[36]: count    7928.000000  
      mean     38.833249  
      std     20.989486  
      min      0.000000  
      25%     24.000000  
      50%     38.000000  
      75%     54.000000  
      max     97.000000  
      Name: age, dtype: float64
```

```
[37]: df_income_train.loc[(df_income_train['rez_esc'].isnull() &_  
    ↪ ((df_income_train['age'] > 7) & (df_income_train['age'] < 17)))]['age'].  
    ↪ describe()
```

```
[37]: count      1.0  
      mean     10.0  
      std      NaN  
      min     10.0  
      25%     10.0  
      50%     10.0  
      75%     10.0  
      max     10.0  
      Name: age, dtype: float64
```

```
[38]: df_income_train[(df_income_train['age'] ==10) & df_income_train['rez_esc'].
      ↪isnull()).head()
df_income_train[(df_income_train['Id'] =='ID_f012e4242')].head()
```

```
[38]:          Id      v2a1  hacdor  rooms  hacapo  v14a  refriger  v18q  \
2514  ID_f012e4242  160000.0        0      6        0      1          1      1

      v18q1  r4h1  ...  SQBescolari  SQBage  SQBhogar_total  SQBedjefe  \
2514    1.0    0  ...              0     100              9         121

      SQBhogar_nin  SQBovercrowding  SQBdependency  SQBmeaned  agesq  Target
2514              1              2.25            0.25     182.25     100       4

[1 rows x 143 columns]
```

```
[39]: for df in [df_income_train, df_income_test]:
      df['rez_esc'].fillna(value=0, inplace=True)
df_income_train[['rez_esc']].isnull().sum()
```

```
[39]: rez_esc      0
      dtype: int64
```

```
[41]: # meaneduc column
```

```
[40]: data = df_income_train[df_income_train['meaneduc'].isnull()].head()

      columns=['edjefe','edjefa','instlevel1','instlevel2']
      data[columns][data[columns]['instlevel1']>0].describe()
```

```
[40]:      edjefe  edjefa  instlevel1  instlevel2
count      0.0      0.0          0.0          0.0
mean       NaN      NaN          NaN          NaN
std        NaN      NaN          NaN          NaN
min        NaN      NaN          NaN          NaN
25%        NaN      NaN          NaN          NaN
50%        NaN      NaN          NaN          NaN
75%        NaN      NaN          NaN          NaN
max        NaN      NaN          NaN          NaN
```

```
[41]: for df in [df_income_train, df_income_test]:
      df['meaneduc'].fillna(value=0, inplace=True)
df_income_train[['meaneduc']].isnull().sum()
```

```
[41]: meaneduc      0
      dtype: int64
```

```
[42]: df_income_test[['meaneduc']].isnull().sum()
```

```
[42]: meaneduc      0
      dtype: int64
```

```
[43]: # SQBmeaned Column
```

```
[44]: data = df_income_train[df_income_train['SQBmeaned'].isnull()].head()

      columns=['edjefe','edjefa','instlevel1','instlevel2']
      data[columns][data[columns]['instlevel1']>0].describe()
```

```
[44]:
```

	edjefe	edjefa	instlevel1	instlevel2
count	0.0	0.0	0.0	0.0
mean	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN

```
[45]: for df in [df_income_train, df_income_test]:
      df['SQBmeaned'].fillna(value=0, inplace=True)
      df_income_train[['SQBmeaned']].isnull().sum()
```

```
[45]: SQBmeaned      0
      dtype: int64
```

```
[46]: df_income_test[['SQBmeaned']].isnull().sum()
```

```
[46]: SQBmeaned      0
      dtype: int64
```

```
[47]: null_counts = df_income_train.isnull().sum()
      null_counts[null_counts > 0].sort_values(ascending=False)
```

```
[47]: Series([], dtype: int64)
```

```
[48]: null_counts = df_income_test.isnull().sum()
      null_counts[null_counts > 0].sort_values(ascending=False)
```

```
[48]: Series([], dtype: int64)
```

```
[49]: # Groupby the household and figure out the number of unique values
      all_equal = df_income_train.groupby('idhogar')['Target'].apply(lambda x: x.
      ↪nunique() == 1)

      # Households where targets are not all equal
```

```
not_equal = all_equal[all_equal != True]
print('There are {} households where the family members do not all have the_
↳same target.'.format(len(not_equal)))
```

There are 85 households where the family members do not all have the same target.

```
[50]: df_income_train[df_income_train['idhogar'] == not_equal.index[0]][['idhogar',
↳'parentesco1', 'Target']]
```

```
[50]:      idhogar  parentesco1  Target
7651  0172ab1d9             0        3
7652  0172ab1d9             0        2
7653  0172ab1d9             0        3
7654  0172ab1d9             1        3
7655  0172ab1d9             0        2
```

```
[51]: households_head = df_income_train.groupby('idhogar')['parentesco1'].sum()

# Find households without a head
households_no_head = df_income_train.loc[df_income_train['idhogar'].
↳isin(households_head[households_head == 0].index), :]

print('There are {} households without a head.'.
↳format(households_no_head['idhogar'].nunique()))
```

There are 15 households without a head.

```
[52]: # Find households without a head and where Target value are different
households_no_head_equal = households_no_head.groupby('idhogar')['Target'].
↳apply(lambda x: x.nunique() == 1)
print('{} Households with no head have different Target value.'.
↳format(sum(households_no_head_equal == False)))
```

0 Households with no head have different Target value.

```
[53]: # Iterating through each household
for household in not_equal.index:
    # Find the correct label (for the head of household)
    true_target = int(df_income_train[(df_income_train['idhogar'] == household)
↳& (df_income_train['parentesco1'] == 1.0)]['Target'])

    # Set the correct label for all members in the household
    df_income_train.loc[df_income_train['idhogar'] == household, 'Target'] =
↳true_target
```

```
# Groupby the household and figure out the number of unique values
all_equal = df_income_train.groupby('idhogar')['Target'].apply(lambda x: x.
    ↪nunique() == 1)

# Households where targets are not all equal
not_equal = all_equal[all_equal != True]
print('There are {} households where the family members do not all have the_
    ↪same target.'.format(len(not_equal)))
```


There are 0 households where the family members do not all have the same target.

```
[54]: # 1 = extreme poverty 2 = moderate poverty 3 = vulnerable households 4 = non_
    ↪vulnerable households
target_counts = heads['Target'].value_counts().sort_index()
target_counts
```

```
[54]: 1      222
      2      442
      3      355
      4     1954
      Name: Target, dtype: int64
```

```
[55]: target_counts.plot.bar(figsize = (8, 6),linewidth = 2,edgecolor =_
    ↪'k',title="Target vs Total_Count")
```

```
[55]: <AxesSubplot:title={'center':'Target vs Total_Count'}>
```



[56]: *# extreme poverty is the smallest count in the train dataset. The dataset is*
↪biased.

```
print(df_income_train.shape)
cols=['SQBescolari', 'SQBage', 'SQBhogar_total', 'SQBedjefe',
      'SQBhogar_nin', 'SQBovercrowding', 'SQBdependency', 'SQBmeaned',
      ↪'agesq']
```

```
for df in [df_income_train, df_income_test]:
    df.drop(columns = cols,inplace=True)
```

```
print(df_income_train.shape)
```

(9557, 143)

(9557, 134)

[57]: `id_ = ['Id', 'idhogar', 'Target']`

```
ind_bool = ['v18q', 'dis', 'male', 'female', 'estadocivil1', 'estadocivil2',
            ↪'estadocivil3',
            'estadocivil4', 'estadocivil5', 'estadocivil6', 'estadocivil7',
```

```

        'parentesco1', 'parentesco2', 'parentesco3', 'parentesco4', □
    ↪ 'parentesco5',
        'parentesco6', 'parentesco7', 'parentesco8', 'parentesco9', □
    ↪ 'parentesco10',
        'parentesco11', 'parentesco12', 'instlevel1', 'instlevel2', □
    ↪ 'instlevel3',
        'instlevel4', 'instlevel5', 'instlevel6', 'instlevel7', □
    ↪ 'instlevel8',
        'instlevel9', 'mobilephone']

ind_ordered = ['rez_esc', 'escolari', 'age']

hh_bool = ['hacdor', 'hacapo', 'v14a', 'refrig', 'paredblolad', 'paredzocalo',
            'paredpreb', 'pisocemento', 'pareddes', 'paredmad',
            'paredzinc', 'paredfibras', 'paredother', 'pisomoscer', 'pisooother',
            'pisonatur', 'pisonotiene', 'pisomadera',
            'techozinc', 'techoentrepiso', 'techocane', 'techootro', 'cielorazo',
            'abastaguadentro', 'abastaguafuera', 'abastaguano',
            'public', 'planpri', 'noelec', 'coopele', 'sanitario1',
            'sanitario2', 'sanitario3', 'sanitario5', 'sanitario6',
            'energcocinar1', 'energcocinar2', 'energcocinar3', 'energcocinar4',
            'elimbasu1', 'elimbasu2', 'elimbasu3', 'elimbasu4',
            'elimbasu5', 'elimbasu6', 'epared1', 'epared2', 'epared3',
            'etecho1', 'etecho2', 'etecho3', 'eviv1', 'eviv2', 'eviv3',
            'tipovivi1', 'tipovivi2', 'tipovivi3', 'tipovivi4', 'tipovivi5',
            'computer', 'television', 'lugar1', 'lugar2', 'lugar3',
            'lugar4', 'lugar5', 'lugar6', 'area1', 'area2']

hh_ordered = [ 'rooms', 'r4h1', 'r4h2', 'r4h3', 'r4m1', 'r4m2', 'r4m3', 'r4t1', □
    ↪ 'r4t2',
        'r4t3', 'v18q1', 'tamhog', 'tamviv', 'hhsize', 'hogar_nin',
        'hogar_adul', 'hogar_mayor', 'hogar_total', 'bedrooms', □
    ↪ 'qmobilephone']

hh_cont = ['v2a1', 'dependency', 'edjefe', 'edjefa', 'meaneduc', 'overcrowding']

```

```

[58]: #Check for redundant household variables
heads = df_income_train.loc[df_income_train['parentesco1'] == 1, :]
heads = heads[id_ + hh_bool + hh_cont + hh_ordered]
heads.shape

```

[58]: (2973, 98)

```

[59]: # Create correlation matrix
corr_matrix = heads.corr()

```



```
# Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(np.
    ↪bool))

# Find index of feature columns with correlation greater than 0.95
to_drop = [column for column in upper.columns if any(abs(upper[column]) > 0.95)]

to_drop
```


```
[59]: ['coopele', 'area2', 'tamhog', 'hhsz', 'hogar_total']
```

```
[60]: corr_matrix.loc[corr_matrix['tamhog'].abs() > 0.9, corr_matrix['tamhog'].abs()
    ↪> 0.9]
```

```
[60]:
```

	r4t3	tamhog	tamviv	hhsz	hogar_total
r4t3	1.000000	0.996884	0.929237	0.996884	0.996884
tamhog	0.996884	1.000000	0.926667	1.000000	1.000000
tamviv	0.929237	0.926667	1.000000	0.926667	0.926667
hhsz	0.996884	1.000000	0.926667	1.000000	1.000000
hogar_total	0.996884	1.000000	0.926667	1.000000	1.000000

```
[61]: sns.heatmap(corr_matrix.loc[corr_matrix['tamhog'].abs() > 0.9,
    ↪corr_matrix['tamhog'].abs() > 0.9],
    annot=True, cmap = plt.cm.flag, fmt='.3f');
```



```
[62]: cols=['tamhog', 'hogar_total', 'r4t3']
for df in [df_income_train, df_income_test]:
    df.drop(columns = cols,inplace=True)

df_income_train.shape
```

[62]: (9557, 131)

```
[63]: #Check for redundant Individual variables
ind = df_income_train[id_ + ind_bool + ind_ordered]
ind.shape
```

[63]: (9557, 39)

```
[64]: # Create correlation matrix
corr_matrix = ind.corr()

# Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(np.
    ↪bool))

# Find index of feature columns with correlation greater than 0.95
to_drop = [column for column in upper.columns if any(abs(upper[column]) > 0.95)]

to_drop
```

[64]: ['female']

```
[65]: # This is simply the opposite of male! We can remove the male flag.
for df in [df_income_train, df_income_test]:
    df.drop(columns = 'male',inplace=True)

df_income_train.shape
```

[65]: (9557, 130)

```
[66]: #lets check area1 and area2 also
# area1, =1 zona urbana
# area2, =2 zona rural
#area2 redundant because we have a column indicating if the house is in a urban_
    ↪zone

for df in [df_income_train, df_income_test]:
    df.drop(columns = 'area2',inplace=True)

df_income_train.shape
```

[66]: (9557, 129)

```
[67]: cols=['Id','idhogar']
for df in [df_income_train, df_income_test]:
    df.drop(columns = cols,inplace=True)

df_income_train.shape
```

[67]: (9557, 127)

```
[68]: df_income_train['Target'].isnull().any().sum()
```

[68]: 0

```
[69]: x_features=df_income_train.iloc[:,0:-1]
y_features=df_income_train.iloc[:,-1]
print(x_features.shape)
print(y_features.shape)
```

(9557, 126)

(9557,)

```
[70]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import
    accuracy_score, confusion_matrix, f1_score, classification_report

x_train,x_test,y_train,y_test=train_test_split(x_features,y_features,test_size=0.
    2,random_state=1)
rmclassifier = RandomForestClassifier()
```

```
[71]: rmclassifier.fit(x_train,y_train)
```

```
[71]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
    criterion='gini', max_depth=None, max_features='auto',
    max_leaf_nodes=None, max_samples=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    n_jobs=None, oob_score=False, random_state=None,
    verbose=0, warm_start=False)
```

```
[72]: y_predict = rmclassifier.predict(x_test)
```

```
[73]: print(accuracy_score(y_test,y_predict))
```

0.946652719665272

```
[74]: print(confusion_matrix(y_test,y_predict))
```

```
[[ 135    0    0   22]
 [   0  285    1   31]
 [   0    2  188   43]
 [   0    2    1 1202]]
```

```
[75]: print(classification_report(y_test,y_predict))
```

	precision	recall	f1-score	support
1	1.00	0.86	0.92	157
2	0.99	0.90	0.94	317
3	0.99	0.81	0.89	233
4	0.93	1.00	0.96	1205
accuracy			0.95	1912
macro avg	0.98	0.89	0.93	1912
weighted avg	0.95	0.95	0.95	1912

```
[76]: y_predict_testdata = rmclassifier.predict(df_income_test)
```

```
[77]: y_predict_testdata
```

```
[77]: array([4, 4, 4, ..., 4, 4, 4])
```

```
[78]: # Predict the accuracy using random forest classifier.
      # Check the accuracy using random forest with cross validation
      from sklearn.model_selection import KFold,cross_val_score
```

```
[79]: seed=7
      kfold=KFold(n_splits=5,random_state=seed,shuffle=True)

      rmclassifier=RandomForestClassifier(random_state=10,n_jobs = -1)
      print(cross_val_score(rmclassifier,x_features,y_features,cv=kfold,scoring='accuracy'))
      results=cross_val_score(rmclassifier,x_features,y_features,cv=kfold,scoring='accuracy')
      print(results.mean()*100)
```

```
[0.94246862 0.94979079 0.94557823 0.94243851 0.94976452]
94.60081361157272
```

```
[80]: num_trees= 100
      rmclassifier=RandomForestClassifier(n_estimators=100, random_state=10,n_jobs = -1)
      print(cross_val_score(rmclassifier,x_features,y_features,cv=kfold,scoring='accuracy'))
      results=cross_val_score(rmclassifier,x_features,y_features,cv=kfold,scoring='accuracy')
```

```
print(results.mean()*100)
```

```
[0.94246862 0.94979079 0.94557823 0.94243851 0.94976452]  
94.60081361157272
```

```
[81]: rmclassifier.fit(x_features,y_features)  
labels = list(x_features)  
feature_importances = pd.DataFrame({'feature': labels, 'importance':  
    ↪rmclassifier.feature_importances_})  
feature_importances=feature_importances[feature_importances.importance>0.015]  
feature_importances.head()
```

```
[81]:   feature  importance  
0     v2a1    0.018653  
2    rooms    0.025719  
9     r4h2    0.020706  
10    r4h3    0.019808  
11    r4m1    0.015271
```

```
[ ]: # Output According to the actions need to be Performed  
  
# Identify the output variable. - Target Variable  
  
# Understand the type of data. - 3 types of Data - Int,Float,Object  
  
# Check if there are any biases in your dataset. - Yes bias was there  
  
# Check whether all members of the house have the same poverty level. - No, It  
    ↪was differing  
  
# Check if there is a house without a family head. - Yes there were some house  
    ↪very less number compared to the total no.of houses  
  
# Set poverty level of the members and the head of the house within a family.-  
    ↪Yes,It has been set on the basis of target varaible  
  
# Count how many null values are existing in columns.  
  
# Remove null value rows of the target variable.  
  
# Predict the accuracy using random forest classifier. - 0.946652719665272 = 0.  
    ↪95 According to classification report  
  
# Check the accuracy using random forest with cross validation. - 0.94246862
```